Quarkus Cheat-Sheet

What is Quarkus?

Quarkus is a Kubernetes Native Java stack tailored for GraalVM & OpenJDK HotSpot, crafted from the best of breed Java libraries and standards. Also focused on developer experience, making things just work with little to no configuration and allowing to do live coding.

Cheat-sheet tested with **Quarkus 0.20.0**.

Getting Started

Quarkus comes with a Maven archetype to scaffold a very simple starting project.

```
mvn io.quarkus:quarkus-maven-plugin:0.20.0:create \
    -DprojectGroupId=org.acme \
    -DprojectArtifactId=getting-started \
    -DclassName="org.acme.quickstart.GreetingResource" \
    -Dpath="/hello"
```

This creates a simple JAX-RS resource called `GreetingResource`.

```
@Path("/hello")
public class GreetingResource {

    @GET
    @Produces(MediaType.TEXT_PLAIN)
    public String hello() {
        return "hello";
    }
}
```

Extensions

Quarkus comes with extensions to integrate with some libraries such as JSON-B, Camel or MicroProfile spec. To list all available extensions just run:

```
./mvnw quarkus:list-extensions
```

Tip You can use `-DsearchPattern=panache` to filter out all extensions except the ones matching the expression.

And to register the extensions into build tool:

```
./mvnw quarkus:add-extension -Dextensions=""
```

Tip `extensions` property supports CSV format to register more than one extension at once.

Application Lifecycle

You can be notified when the application starts/stops by observing `StartupEvent` and `ShutdownEvent` events.

```
@ApplicationScoped
public class ApplicationLifecycle {
    void onStart(@Observes StartupEvent event) {}
    void onStop(@Observes ShutdownEvent event) {}
}
```

Adding Configuration Parameters

To add configuration to your application, Quarkus relies on MicroProfile Config spec.

```
@ConfigProperty(name = "greetings.message")
String message;

@ConfigProperty(name = "greetings.message",
                defaultValue = "Hello")
String messageWithDefault;

@ConfigProperty(name = "greetings.message")
Optional<String> optionalMessage;
```

Properties can be set as:

- Environment variables (`GREETINGS_MESSAGE`).

- System properties (`-Dgreetings.message`).

- Resources `src/main/resources/application.properties`.

- External config directory under the current working directory: `config/application.properties`.

```
greetings.message = Hello World
```

Tip `Array`, `List` and `Set` are supported. The delimiter is comma (`,`) char and `\` is the escape char.

Configuration Profiles

Quarkus allow you to have multiple configuration in the same file (`application.properties`).

The syntax for this is `%{profile}.config.key=value`.

```
quarkus.http.port=9090
%dev.quarkus.http.port=8181
```

HTTP port will be 9090, unless the 'dev' profile is active.

Default profiles are:

- `dev`: Activated when in development mode (`quarkus:dev`).

- `test`: Activated when running tests.

- `prod`: The default profile when not running in development or test mode

You can create custom profile names by enabling the profile either setting `quarkus.profile` system property or `QUARKUS_PROFILE` environment variable.

```
quarkus.http.port=9090
%staging.quarkus.http.port=9999
```

And enable it `quarkus.profile=staging`.

You can also set it in the build tool:

```
<groupId>org.apache.maven.plugins</groupId>
<artifactId>maven-surefire-plugin</artifactId>
<version>${surefire-plugin.version}</version>
<configuration>
    <systemPropertyVariables>
        <quarkus.test.profile>foo</quarkus.test.profile>
        <buildDirectory>${project.build.directory}</buildDirectory>
    </systemPropertyVariables>
</configuration>
```

Tip Same for `maven-failsafe-plugin`.

```
test {
    useJUnitPlatform()
    systemProperty "quarkus.test.profile", "foo"
}
```

Custom Loader

You can implement your own `ConfigSource` to load configuration from different places than the default ones provided by Quarkus. For example, database, custom XML, REST Endpoints, ...

You need to create a new class and implement `ConfigSource` interface:

```java
package com.acme.config;
public class InMemoryConfig implements ConfigSource {

    private Map<String, String> prop = new HashMap<>();

    public InMemoryConfig() {
        // Init properties
    }

    @Override
    public int getOrdinal() {
        // The highest ordinal takes precedence
        return 900;
    }

    @Override
    public Map<String, String> getProperties() {
        return prop;
    }

    @Override
    public String getValue(String propertyName) {
        return prop.get(propertyName);
    }

    @Override
    public String getName() {
        return "MemoryConfigSource";
    }
}
```

Then you need to register the `ConfigSource` as Java service. Create a file /META-INF/services/org.eclipse.microprofile.config.spi.ConfigSource with next content:

```
com.acme.config.InMemoryConfig
```

Custom Converters

You can implement your own conversion types from String. Implement org.eclipse.microprofile.config.spi.Converter interface:

```java
@Priority(DEFAULT_QUARKUS_CONVERTER_PRIORITY + 100)
public class CustomInstantConverter
    implements Converter<Instant> {

    @Override
    public Instant convert(String value) {
        if ("now".equals(value.trim())) {
            return Instant.now();
        }
        return Instant.parse(value);
    }
}
```

`@Priority` annotation is used to override the default `InstantConverter`.

Then you need to register the `Converter` as Java service. Create a file /META-INF/services/org.eclipse.microprofile.config.spi.Converter with next content:

```
com.acme.config.CustomInstantConverter
```

Custom Context Path

By default Undertow will serve content from under the root context. If you want to change this you can use the `quarkus.servlet.context-path` config key to set the context path.

Injection

Quarkus is based on CDI 2.0 to implement injection of code. It is not fully supported and only a subset of the specification is implemented.

```java
@ApplicationScoped
public class GreetingService {

    public String message(String message) {
        return message.toUpperCase();
    }
}
```

Scope annotation is mandatory to make the bean discoverable.

```java
@Inject
GreetingService greetingService;
```

Important Quarkus is designed with Substrate VM in mind. For this reason, we encourage you to use *package-private* scope instead of *private*.

Produces You can also create a factory of an object by using `@javax.enterprise.inject.Produces` annotation.

```java
@Produces
@ApplicationScoped
Message message() {
    Message m = new Message();
    m.setMsn("Hello");
    return m;
}

@Inject
Message msg;
```

You can use qualifiers to return different implementations of the same interface or to customize the configuration of the bean.

```java
@Qualifier
@Retention(RUNTIME)
@Target({TYPE, METHOD, FIELD, PARAMETER})
public @interface Quote {
    @Nonbinding String value();
}

@Produces
@Quote("")
Message message(InjectionPoint msg) {
    Message m = new Message();
    m.setMsn(
        msg.getAnnotated()
        .getAnnotation(Quote.class)
        .value()
    );

    return m;
}

@Inject
@Quote("Aloha Beach")
Message message;
```

JSON Marshalling/Unmarshalling

To work with `JSON-B` you need to add a dependency:

```
./mvnw quarkus:add-extension
    -Dextensions="io.quarkus:quarkus-resteasy-jsonb"
```

Any POJO is marshaled/unmarshalled automatically.

```java
public class Sauce {
    private String name;
    private long scovilleHeatUnits;

    // getter/setters
}
```

JSON equivalent:

```json
{
    "name":"Blair's Ultra Death",
    "scovilleHeatUnits": 1100000
}
```

In a `POST` endpoint example:

```
@POST
@Consumes(MediaType.APPLICATION_JSON)
public Response create(Sauce sauce) {
    // Create Sauce
    return Response.created(URI.create(sauce.getId()))
            .build();
}
```

To work with `Jackson` you need to add:

```
./mvnw quarkus:add-extension
    -Dextensions="quarkus-resteasy-jackson"
```

If you don't want to use the default `ObjectMapper` you can customize it by:

```
@ApplicationScoped
public class CustomObjectMapperConfig {
    @Singleton
    @Produces
    public ObjectMapper objectMapper() {
        ObjectMapper objectMapper = new ObjectMapper();
        // perform configuration
        return objectMapper;
    }
}
```

Validator

Quarkus uses **Hibernate Validator** to validate input/output of REST services and business services using Bean validation spec.

```
./mvnw quarkus:add-extension
    -Dextensions="io.quarkus:quarkus-hibernate-validator"
```

Annotate POJO objects with validator annotations such as: `@NotNull`, `@Digits`, `@NotBlank`, `@Min`, `@Max`, …

```
public class Sauce {

    @NotBlank(message = "Name may not be blank")
    private String name;
    @Min(0)
    private long scovilleHeatUnits;

    // getter/setters
}
```

To validate an object use `@Valid` annotation:

```
public Response create(@Valid Sauce sauce) {}
```

Tip If a validation error is triggered, a violation report is generated and serialized as JSON. If you want to manipulate the output, you need to catch in the code the `ConstraintViolationException` exception.

Create Your Custom Constraints

First you need to create the custom annotation:

```
@Target({ METHOD, FIELD, ANNOTATION_TYPE, CONSTRUCTOR,
          PARAMETER, TYPE_USE })
@Retention(RUNTIME)
@Documented
@Constraint(validatedBy = { NotExpiredValidator.class})
public @interface NotExpired {

    String message() default "Sauce must not be expired";
    Class<?>[] groups() default { };
    Class<? extends Payload>[] payload() default { };

}
```

You need to implement the validator logic in a class that implements `ConstraintValidator`.

```
public class NotExpiredValidator
    implements ConstraintValidator<NotExpired, LocalDate>
    {

    @Override
    public boolean isValid(LocalDate value,
                    ConstraintValidatorContext ctx) {
        if ( value == null ) return true;
        LocalDate today = LocalDate.now();
        return ChronoUnit.YEARS.between(today, value) > 0;
    }
}
```

And use it normally:

```
@NotExpired
@JsonbDateFormat(value = "yyyy-MM-dd")
private LocalDate expired;
```

Manual Validation

You can call the validation process manually instead of relaying to `@Valid` by injecting `Validator` class.

```
@Inject
Validator validator;
```

And use it:

```
Set<ConstraintViolation<Sauce>> violations =
        validator.validate(sauce);
```

Logging

You can configure how Quarkus logs:

```
quarkus.log.console.enable=true
quarkus.log.console.level=DEBUG
quarkus.log.console.color=false
quarkus.log.category."com.lordofthejars".level=DEBUG
```

Prefix is `quarkus.log.`

Property	Default	Description
console.enable	true	Console logging enabled.
console.format	%d{yyyy-MM-dd HH:mm:ss,SSS} %-5p [%c{3.}] (%t) %s%e%n	Format pattern to use for logging.
console.level	INFO	Minimum log level.
console.color	INFO	Allow color rendering.
file.enable	false	File logging enabled.
file.format	%d{yyyy-MM-dd HH:mm:ss,SSS} %h %N[%i] %-5p [%c{3.}] (%t) %s%e%n	Format pattern to use for logging.
file.level	ALL	Minimum log level.
file.path	quarkus.log	The path to log file.
file.rotation.max-file-size		The maximum file size of the log file.

Property	Default	Description
file.rotation.max-backup-index	1	The maximum number of backups to keep.
file.rotation.file-suffix		Rotating log file suffix.
file.rotation.rotate-on-boot	true	Indicates rotate logs at bootup.
file.async	false	Log asynchronously.
file.async.queue-length	512	The queue length to use before flushing writing.
file.async.overflow	BLOCK	Action when queue is full.
syslog.enable	false	syslog logging is enabled.
syslog.format	`%d{yyyy-MM-dd HH:mm:ss,SSS} %h %N[%i] %-5p [%c{3.}] (%t) %s%e%n`	The format pattern to use for logging to syslog.
syslog.level	ALL	The minimum log level to write to syslog.
syslog.endpoint	localhost:514	The IP address and port of the syslog server.
syslog.app-name	Current process name.	The app name used when formatting the message in RFC5424 format.
syslog.hostname	Current hostname.	The name of the host the messages are being sent from.

Property	Default	Description
syslog.facility	USER_LEVEL	Priority of the message as defined by RFC-5424 and RFC-3164.
syslog.syslog-type	RFC5424	The syslog type of format message.
syslog.protocol	TCP	Protocol used.
syslog.use-counting-framing	false	Message prefixed with the size of the message.
syslog.truncate	true	Message should be truncated.
syslog.block-on-reconnect	true	Block when attempting to reconnect.
syslog.async	false	Log asynchronously.
syslog.async.queue-length	512	The queue length to use before flushing writing.
syslog.async.overflow	BLOCK	Action when queue is full.
category."<category-name>".level	INFO	Minimum level category.
level	INFO	Default minimum level.

Rest Client

Quarkus implements MicroProfile Rest Client spec:

```
./mvnw quarkus:add-extension
  -Dextensions="io.quarkus:quarkus-rest-client"
```

To get content from http://worldclockapi.com/api/json/cet/now you need to create a service interface:

```
@Path("/api")
@RegisterRestClient
public interface WorldClockService {

    @GET @Path("/json/cet/now")
    @Produces(MediaType.APPLICATION_JSON)
    WorldClock getNow();

    @GET
    @Path("/json/{where}/now")
    @Produces(MediaType.APPLICATION_JSON)
    WorldClock getSauce(@BeanParam
            WorldClockOptions worldClockOptions);

}
```

```
public class WorldClockOptions {
    @HeaderParam("Authorization")
    String auth;

    @PathParam("where")
    String where;
}
```

And configure the hostname at `application.properties`:

```
org.acme.quickstart.WorldClockService/mp-rest/url=
        http://worldclockapi.com
```

Injecting the client:

```
@Inject
@RestClient
WorldClockService worldClockService;
```

If invokation happens within JAX-RS, you can propagate headers from incoming to outgoing by using next property.

```
org.eclipse.microprofile.rest.client.propagateHeaders=
        Authorization,MyCustomHeader
```

Tip You can still use the JAX-RS client without any problem `ClientBuilder.newClient().target(…)`

Adding headers

You can customize the headers passed by implementing MicroProfile `ClientHeadersFactory` annotation:

```
@RegisterForReflection
public class BaggageHeadersFactory
            implements ClientHeadersFactory {
    @Override
    public MultivaluedMap<String, String> update(
        MultivaluedMap<String, String> incomingHeaders,
        MultivaluedMap<String, String> outgoingHeaders) {}
}
```

And registering it in the client using `RegisterClientHeaders` annotation.

```
@RegisterClientHeaders(BaggageHeadersFactory.class)
@RegisterRestClient
public interface WorldClockService {}
```

Or statically set:

```
@GET
@ClientHeaderParam(name="X-Log-Level", value="ERROR")
Response getNow();
```

Asynchronous

A method on client interface can return a `CompletionStage` class to be executed asynchronously.

```
@GET @Path("/json/cet/now")
@Produces(MediaType.APPLICATION_JSON)
CompletionStage<WorldClock> getNow();
```

Testing

Quarkus archetype adds test dependencies with JUnit 5 and Rest-Assured library to test REST endpoints.

```
@QuarkusTest
public class GreetingResourceTest {

    @Test
    public void testHelloEndpoint() {
        given()
          .when().get("/hello")
          .then()
            .statusCode(200)
            .body(is("hello"));
    }
}
```

Test port can be set in `quarkus.http.test-port` property.

You can also inject the URL where Quarkus is started:

```
@TestHTTPResource("index.html")
URL url;
```

Quarkus Test Resource

You can execute some logic before the first test run (`start`) and execute some logic at the end of the test suite (`stop`).

You need to create a class implementing `QuarkusTestResourceLifecycleManager` interface and register it in the test via `@QuarkusTestResource` annotation.

```
public class MyCustomTestResource
    implements QuarkusTestResourceLifecycleManager {

    @Override
    public Map<String, String> start() {
        // return system properties that
        // should be set for the running test
        return Collections.emptyMap();
    }

    @Override
    public void stop() {
    }

    // optional
    @Override
    public void inject(Object testInstance) {
    }

    // optional
    @Override
    public int order() {
        return 0;
    }
}
```

Important Returning new system properties implies running parallel tests in different JVMs.

And the usage:

```
@QuarkusTestResource(MyCustomTestResource.class)
public class MyTest {
}
```

Mocking

If you need to provide an alternative implementation of a service (for testing purposes) you can do it by using CDI `@Alternative` annotation using it in the test service placed at `src/test/java`:

```
@Alternative
@Priority(1)
@ApplicationScoped
public class MockExternalService extends ExternalService {}
```

Important This does not work when using native image testing.

A stereotype annotation `io.quarkus.test.Mock` is provided declaring `@Alternative`, `@Priority(1)` and `@Dependent`.

Interceptors

Tests are actually full CDI beans, so you can apply CDI interceptors:

```
@QuarkusTest
@Stereotype
@Transactional
@Retention(RetentionPolicy.RUNTIME)
@Target(ElementType.TYPE)
public @interface TransactionalQuarkusTest {
}

@TransactionalQuarkusTest
public class TestStereotypeTestCase {}
```

Test Coverage Due the nature of Quarkus to calculate correctly the coverage information with JaCoCo, you might need offline instrumentation. I recommend reading **this document** to understand how JaCoCo and Quarkus works and how you can configure JaCoCo to get correct data.

Native Testing

To test native executables annotate the test with `@SubstrateTest`.

Persistence

Quarkus works with JPA(Hibernate) as persistence solution. But also provides an **Active Record pattern** implementation under Panache project.

To use database access you need to add Quarkus JDBC drivers instead of the original ones. At this time `H2`, `MariaDB`, `MSSQL` and `PostgreSQL` drivers are supported.

```
./mvnw quarkus:add-extension
  -Dextensions="io.quarkus:quarkus-hibernate-orm-panache,
            io.quarkus:quarkus-jdbc-mariadb"
```

```
@Entity
public class Developer extends PanacheEntity {

    // id field is implicit

    public String name;
}
```

And configuration in `src/main/resources/application.properties`:

```
quarkus.datasource.url=jdbc:mariadb://localhost:3306/mydb
quarkus.datasource.driver=org.mariadb.jdbc.Driver
quarkus.datasource.username=developer
quarkus.datasource.password=developer
quarkus.hibernate-orm.database.generation=update
```

List of datasource parameters.

`quarkus.datasource` as prefix is skipped in the next table.

Parameter	Type
driver	String
url	String
username	String
password	String
min-size	Integer
max-size	Integer
initial-size	Integer

Parameter	Type
background-validation-interval	java.time.Duration
acquisition-timeout	java.time.Duration
leak-detection-interval	java.time.Duration
idle-removal-interval	java.time.Duration
transaction-isolation-level	io.quarkus.agroal.runtime.TransactionIsolationLevel
enable-metrics	Boolean
xa	Boolean

Hibernate configuration properties. Prefix `quarkus.hibernate-orm` is skipped.

Parameter	Description	Values[Default]
dialect	Class name of the Hibernate ORM dialect.	Not necessary to set.
dialect.storage-engine	The storage engine when the dialect supports multiple storage engines.	Not necessary to set.

Parameter	Description	Values[Default]
sql-load-script	Name of the file containing the SQL statements to execute when starts. `no-file` force Hibernate to skip SQL import.	`import.sql`
batch-fetch-size	The size of the batches.	-1 disabled.
query.query-plan-cache-max-size	The maximum size of the query plan cache.	
query.default-null-ordering	Default precedence of null values in `ORDER BY`. none, first, last.	[none], first, last.
database.generation	Database schema is generation.	[none], create, drop-and-create, drop, update.
database.generation.halt-on-error	Stop on the first error when applying the schema.	[flase], true
database.default-catalog	Default catalog.	
database.default-schema	Default Schema.	
database.charset	Charset.	
jdbc.timezone	Time Zone JDBC driver.	

Parameter	Description	Values[Default]
`jdbc.statement-fetch-size`	Number of rows fetched at a time.	
`jdbc.statement-batch-size`	Number of updates sent at a time.	
`log.sql`	Show SQL logs	[`false`], `true`
`log.jdbc-warnings`	Collect and show JDBC warnings.	[`false`], `true`
`statistics`	Enable statiscs collection.	[`false`], `true`

Database operations:

```java
// Insert
Developer developer = new Developer();
developer.name = "Alex";
developer.persist();

// Find All
Developer.findAll().list();

// Find By Query
Developer.find("name", "Alex").firstResult();

// Delete
Developer developer = new Developer();
developer.id = 1;
developer.delete();

// Delete By Query
long numberOfDeleted = Developer.delete("name", "Alex");
```

Remember to annotate methods with `@Transactional` annotation to make changes persisted in the database.

If queries start with the keyword `from` then they are treated as *HQL* query, if not then next short form is supported:

- `order by` which expands to `from EntityName order by` …
- `<columnName>` which expands to `from EntityName where <columnName>=?`
- `<query>` which is expanded to `from EntityName where <query>`

Static Methods

Field	Parameters	Return
`findById`	`Object`	Returns object or null if not found.
`find`	`String, [Object…, Map<String, Object>, Parameters]`	Lists of entities meeting given query with parameters set.
`find`	`String, Sort, [Object…, Map<String, Object>, Parameters]`	Lists of entities meeting given query with parameters set sorted by `Sort` attribute/s.
`findAll`		Finds all entities.
`findAll`	`Sort`	Finds all entities sorted by `Sort` attribute/s.
`stream`	`String, [Object…, Map<String, Object>, Parameters]`	`java.util.stream.Stream` of entities meeting given query with parameters set.
`stream`	`String, Sort, [Object…, Map<String, Object>, Parameters]`	`java.util.stream.Stream` of entities meeting given query with parameters set sorted by `Sort` attribute/s.
`streamAll`		`java.util.stream.Stream` of all entities.
`streamAll`	`Sort`	`java.util.stream.Stream` of all entities sorted by `Sort` attribute/s.
`count`		`Number of entities.
`count`	`String, [Object…, Map<String, Object>, Parameters]`	Number of entities meeting given query with parameters set.
`deleteAll`		Number of deleted entities.
`delete`	`String, [Object…, Map<String, Object>, Parameters]`	Number of deleted entities meeting given query with parameters set.
`persist`	`[Iterable, Steram, Object…]`	

If entities are defined in external JAR, you need to enable in these projects the `Jandex` plugin in project.

```xml
<plugin>
    <groupId>org.jboss.jandex</groupId>
    <artifactId>jandex-maven-plugin</artifactId>
    <version>1.0.3</version>
    <executions>
        <execution>
            <id>make-index</id>
            <goals>
                <goal>jandex</goal>
            </goals>
        </execution>
    </executions>
    <dependencies>
        <dependency>
            <groupId>org.jboss</groupId>
            <artifactId>jandex</artifactId>
            <version>2.1.1.Final</version>
        </dependency>
    </dependencies>
</plugin>
```

DAO pattern

Also supports *DAO* pattern with `PanacheRepository<TYPE>`.

```java
@ApplicationScoped
public class DeveloperRepository
    implements PanacheRepository<Person> {
    public Person findByName(String name){
        return find("name", name).firstResult();
    }
}
```

EntityManager You can inject `EntityManager` in your classes:

```java
@Inject
EntityManager em;
```

Flushing

You can force flush operation by calling `.flush()` or `.persistAndFlush()` to make it in a single call.

Important This flush is less efficient and you still need to commit transaction.

Testing

There is a Quarkus Test Resource that starts and stops H2 server before and after test suite.

Register next dependency `io.quarkus:quarkus-test-h2:test`.

And annotate the test:

```
@QuarkusTestResource(H2DatabaseTestResource.class)
public class FlywayTestResources {
}
```

Transactions The easiest way to define your transaction boundaries is to use the `@Transactional` annotation.

Transactions are mandatory in case of none idempotent operations.

```
@Transactional
public void createDeveloper() {}
```

You can control the transaction scope:

- `@Transactional(REQUIRED)` (default): starts a transaction if none was started, stays with the existing one otherwise.

- `@Transactional(REQUIRES_NEW)`: starts a transaction if none was started; if an existing one was started, suspends it and starts a new one for the boundary of that method.

- `@Transactional(MANDATORY)`: fails if no transaction was started ; works within the existing transaction otherwise.

- `@Transactional(SUPPORTS)`: if a transaction was started, joins it ; otherwise works with no transaction.

- `@Transactional(NOT_SUPPORTED)`: if a transaction was started, suspends it and works with no transaction for the boundary of the method; otherwise works with no transaction.

- `@Transactional(NEVER)`: if a transaction was started, raises an exception; otherwise works with no transaction.

You can configure the default transaction timeout using `quarkus.transaction-manager.default-transaction-timeout` configuration property. By default it is set to 60 seconds.

You can set a timeout property, in seconds, that applies to transactions created within the annotated method by using `@TransactionConfiguration` annotation.

```
@Transactional
@TransactionConfiguration(timeout=40)
public void createDeveloper() {}
```

If you want more control over transactions you can inject `UserTransaction` and use a programmatic way.

```
@Inject UserTransaction transaction

transaction.begin();
transaction.commit();
transaction.rollback();
```

Flyway

Quarkus integrates with **Flyway** to help you on database schema migrations.

```
./mvnw quarkus:add-extension
  -Dextensions="io.quarkus:quarkus-flyway"
```

Then place migration files to the migrations folder (`classpath:db/migration`).

You can inject `org.flywaydb.core.Flyway` to programmatically execute the migration.

```
@Inject
Flyway flyway;

flyway.migrate();
```

Or can be automatically executed by setting `migrate-at-start` property to `true`.

```
quarkus.flyway.migrate-at-start=true
```

List of Flyway parameters.

`quarkus.` as prefix is skipped in the next table.

Parameter	Default	Description
flyway.migrate-at-start	false	Flyway migration automatically.
flyway.locations	classpath:db/migration	CSV locations to scan recursively for migrations. Supported prefixes `classpath` and `filesystem`
flyway.connect-retries	0	The maximum number of retries when attempting to connect.
flyway.schemas	none	CSV case-sensitive list of schemas managed.
flyway.table	flyway_schema_history	The name of Flyway's schema history table.
flyway.sql-migration-prefix	V	Prefix for versioned SQL migrations.
flyway.repeatable-sql-migration-prefix	R	Prefix for repeatable SQL migrations.
flyway.baseline-on-migrate		Only migrations above **baseline-version** will then be applied.
flyway.baseline-version	1	Version to tag an existing schema with when executing baseline.

Parameter	Default	Description
flyway.baseline-description	Flyway Baseline	Description to tag an existing schema with when executing baseline

Hibernate Search

Quarkus integrates with Elasticsearch to provide a full-featured full-text search using Hibernate Search API.

```
./mvnw quarkus:add-extension
    -Dextensions="quarkus-hibernate-search-elasticsearch"
```

You need to annotate your model with Hibernate Search API to index it:

```
@Entity
@Indexed
public class Author extends PanacheEntity {

    @FullTextField(analyzer = "english")
    public String bio;

    @FullTextField(analyzer = "name")
    @KeywordField(name = "firstName_sort",
        sortable = Sortable.YES,
        normalizer = "sort")
    public String firstName;

    @OneToMany
    @IndexedEmbedded
    public List<Book> books;

}
```

Important It is not mandatory to use Panache.

You need to define the analyzers and normalizers defined in annotations. You only need to implement `ElasticsearchAnalysisConfigurer` interface and configure it.

```
public class MyQuarkusAnalysisConfigurer
        implements ElasticsearchAnalysisConfigurer {

    @Override
    public void configure(
        ElasticsearchAnalysisDefinitionContainerContext ctx)
    {
        ctx.analyzer("english").custom()
            .withTokenizer("standard")
            .withTokenFilters("asciifolding",
                "lowercase", "porter_stem");

        ctx.normalizer("sort").custom()
            .withTokenFilters("asciifolding", "lowercase");
    }
}
```

Use Hibernate Search in REST service:

```
public class LibraryResource {

    @Inject
    EntityManager em;

    @Transactional
    public List<Author> searchAuthors(
        @QueryParam("pattern") String pattern) {
        return Search.getSearchSession(em)
            .search(Author.class)
            .predicate(f ->
                pattern == null || pattern.isEmpty() ?
                    f.matchAll() :
                    f.simpleQueryString()
                        .onFields("firstName",
                            "lastName", "books.title")
                        .matching(pattern)
            )
            .sort(f -> f.byField("lastName_sort")
            .then().byField("firstName_sort"))
            .fetchHits();
    }
}
```

IMPORTANT If you are importing data without using Hibernate ORM, you need to index data by using `Search.getSearchSession(em).createIndexer().startAndWait()` at startup time.

Configure the extension in `application.properties`:

```
quarkus.hibernate-search.elasticsearch.version=7
quarkus.hibernate-search.elasticsearch.
    analysis-configurer=MyQuarkusAnalysisConfigurer
quarkus.hibernate-search.elasticsearch.
    automatic-indexing.synchronization-strategy=searchable
quarkus.hibernate-search.elasticsearch.
    index-defaults.lifecycle.strategy=drop-and-create
quarkus.hibernate-search.elasticsearch.
    index-defaults.lifecycle.required-status=yellow
```

List of Hibernate-Elasticsearch properties prefixed with `quarkus.hibernate-search.elasticsearch`:

Parameter	Description
backends	Map of configuration of additional backends.
version	Version of Elasticsearch
analysis-configurer	Class or name of the neab used to configure.
hosts	List of Elasticsearch servers hosts.
username	Username for auth.
password	Password for auth.
connection-timeout	Duration of connection timeout.
max-connections	Max number of connections to servers.
max-connections-per-route	Max number of connections to server.
indexes	Per-index specific configuration.
discovery.enabled	Enables automatic discovery.
discovery.refresh-interval	Refresh interval of node list.
discovery.default-scheme	Scheme to be used for the new nodes.
automatic-indexing.synchronization-strategy	Status for which you wait before considering the operation completed (`queued`, `committed` or `searchable`).

Parameter	Description
`automatic-indexing.enable-dirty-check`	When enabled, re-indexing of is skipped if the changes are on properties that are not used when indexing.
`index-defaults.lifecycle.strategy`	Index lifecycle (`none`, `validate`, `update`, `create`, `drop-and-create`, `drop-abd-create-drop`)
`index-defaults.lifecycle.required-status`	Minimal cluster status (`green`, `yellow`, `red`)
`index-defaults.lifecycle.required-status-wait-timeout`	Waiting time before failing the bootstrap.
`index-defaults.refresh-after-write`	Set if index should be refreshed after writes.

Possible annotations:

Parameter	Description
`@Indexed`	Register entity as full text index
`@FullTextField`	Full text search. Need to set an analyzer to split tokens.
`@KeywordField`	The string is kept as one single token but can be normalized.
`IndexedEmbedded`	Include the Book fields into the Author index.
`@ContainerExtraction`	Sets how to extract a value from container, e.g from a `Map`.
`@DocumentId`	Map an unusual entity identifier to a document identifier.
`@GenericField`	Full text index for any supported type.
`@IdentifierBridgeRef`	Reference to the identifier bridge to use for a @DocumentId

Parameter	Description
`@IndexingDependency`	How a dependency of the indexing process to a property should affect automatic reindexing.
`@ObjectPath`	
`@ScaledNumberField`	For `java.math.BigDecimal` or `java.math.BigInteger` that you need higher precision.

Amazon DynamoDB

Quarkus integrates with Amazon DynamoDB:

```
./mvnw quarkus:add-extension
  -Dextensions="quarkus-amazon-dynamodb"
```

```
void onStart(@Observes StartupEvent ev) {
    DynamoDbAsyncClient asyncClient = DynamoDbAsyncClient.create();
    DynamoDbClient client = DynamoDbClient.create();
}
```

Neo4j

Quarkus integrates with Neo4j:

```
./mvnw quarkus:add-extension
  -Dextensions="quarkus-neo4j"
```

```
@Inject
org.neo4j.driver.Driver driver;
```

Configuration properties:

`quarkus.neo4j` as prefix is skipped in the next table.

Parameter	Default	Description
`uri`	`bolt://localhost:7687`	URI of Neo4j.
`authentication.username`	`neo4j`	Username.
`authentication.password`	`neo4j`	Password.
`authentication.disabled`	`false`	Disable

Parameter	Default	Descript
`pool.metrics-enabled`	`false`	Enable metrics.
`pool.log-leaked-sessions`	`false`	Enable sessions logging.
`pool.max-connection-pool-size`	`100`	Max of connecti
`pool.max-connection-lifetime`	`1H`	Pooled connecti older w closed removed the pool
`pool.connection-acquisition-timeout`	`1M`	Timout connecti adquisat
`pool.idle-time-before-connection-test`	`-1`	Pooled connecti idled i pool for than timeout tested they are

As Neo4j uses SSL communication by default, to create a native executable you need to compile with next options GraalVM options:

```
-H:EnableURLProtocols=http,https   --enable-all-
security-services -H:+JNI
```

And Quarkus Maven Plugin with next configuration:

```xml
<artifactId>quarkus-maven-plugin</artifactId>
<executions>
    <execution>
        <id>native-image</id>
        <goals>
            <goal>native-image</goal>
        </goals>
        <configuration>
            <enableHttpUrlHandler>true</enableHttpUrlHandler>
            <enableHttpsUrlHandler>true</enableHttpsUrlHandler
            <enableAllSecurityServices>true</enableAllSecurity:
            <enableJni>true</enableJni>
        </configuration>
    </execution>
</executions>
```

Alternatively, and as a not recommended way in production, you can disable SSL and Quarkus will disable Bolt SSL as well. `quarkus.ssl.native=false`.

If you are using Neo4j 4.0, you can use fully reactive. Add next depenency management `io.projectreactor:reactor-bom:Californium-SR4:pom:import` and dependency: `io.projectreactor:reactor-core`.

```
public Publisher<String> get() {
        return Flux.using(driver::rxSession, ...);
}
```

MongoDB Client

Quarkus integrates with MongoDB:

```
./mvnw quarkus:add-extension
  -Dextensions="quarkus-mongodb-client"
```

```
@Inject
com.mongodb.client.MongoClient client;

@Inject
io.quarkus.mongodb.ReactiveMongoClient client;
```

```
quarkus.mongodb.connection-string=mongodb://localhost:27018
quarkus.mongodb.write-concern.journal=false
```

`quarkus.mongodb` as prefix is skipped in the next table.

Parameter	Type	Description
connection-string	String	MongoDB connection URI.
hosts	List<String>	Addresses passed as host:port.
application-name	String	Application name.
max-pool-size	Int	Maximum number of connections.
min-pool-size	Int	Minimum number of connections.

Parameter	Type	Description
max-connection-idle-time	Duration	Idle time of a pooled connection.
max-connection-life-time	Duration	Life time of pooled connection.
wait-queue-timeout	Duration	Maximum wait time for new connection.
maintenance-frequency	Duration	Time period between runs of maintenance job.
maintenance-initial-delay	Duration	Time to wait before running the first maintenance job.
wait-queue-multiple	Int	Multiplied with `max-pool-size` gives max numer of threads waiting.
connection-timeout	Duration	
socket-timeout	Duration	
tls-insecure	boolean [false]	Insecure TLS.
tls	boolean [false]	Enable TLS
replica-set-name	String	Implies hosts given are a seed list.
server-selection-timeout	Duration	Time to wait for server selection.

Parameter	Type	Description
local-threshold	Duration	Minimum ping time to make a server eligible.
heartbeat-frequency	Duration	Frequency to determine the state of servers.
read-preference	primary, primaryPreferred, secondary, secondaryPreferred, nearest	Read preferences.
max-wait-queue-size	Int	Max number of concurrent operations allowed to wait.
write-concern.safe	boolean [true]	Ensures are writes are ack.
write-concern.journal	boolean [true]	Journal writing aspect.
write-concern.w	String	Value to all write commands.
write-concern.retry-writes	boolean [false]	Retry writes if network fails.
write-concern.w-timeout	Duration	Timeout to all write commands.
credentials.username	String	Username.
credentials.password	String	Password.
credentials.auth-mechanism	MONGO-CR, GSSAPI, PLAIN, MONGODB-X509	
credentials.auth-source	String	Source of the authentication credentials.

Parameter	Type	Description
credentials.auth-mechanism-properties	Map<String, String>	Authentication mechanism properties.

Reactive Programming

Quarkus implements MicroProfile Reactive spec and uses RXJava2 to provide reactive programming model.

```
./mvnw quarkus:add-extension
  -Dextensions="
    io.quarkus:quarkus-smallrye-reactive-streams-operators"
```

Asynchronous HTTP endpoint is implemented by returning Java `CompletionStage`. You can create this class either manually or using MicroProfile Reactive Streams spec:

```
@GET
@Path("/reactive")
@Produces(MediaType.TEXT_PLAIN)
public CompletionStage<String> getHello() {
    return ReactiveStreams.of("h", "e", "l", "l", "o")
    .map(String::toUpperCase)
    .toList()
    .run()
    .thenApply(list -> list.toString());
}
```

Creating streams is also easy, you just need to return `Publisher` object.

```
@GET
@Path("/stream")
@Produces(MediaType.SERVER_SENT_EVENTS)
public Publisher<String> publishers() {
        return Flowable
        .interval(500, TimeUnit.MILLISECONDS)
        .map(s -> atomicInteger.getAndIncrement())
        .map(i -> Integer.toString(i));
}
```

Reactive Messaging

Quarkus relies on MicroProfile Reactive Messaging spec to implement reactive messaging streams.

```
mvn quarkus:add-extension
    -Dextensions="
        io.quarkus:quarkus-smallrye-reactive-messaging"
```

You can just start using in-memory streams by using `@Incoming` to produce data and `@Outgoing` to consume data.

Produce every 5 seconds one piece of data.

```
@ApplicationScoped
public class ProducerData {

    @Outgoing("my-in-memory")
    public Flowable<Integer> generate() {
        return Flowable.interval(5, TimeUnit.SECONDS)
            .map(tick -> random.nextInt(100));
    }
}
```

If you want to dispatch to all subscribers you can annotate the method with `@Broadcast`.

Consumes generated data from `my-in-memory` stream.

```
@ApplicationScoped
public class ConsumerData {
    @Incoming("my-in-memory")
    public void randomNumber(int randomNumber) {
        System.out.println("Received " + randomNumber);
    }
}
```

You can also inject an stream as a field:

```
@Inject
@Stream("my-in-memory") Publisher<Integer> randomRumbers;
```

```
@Inject @Stream("generated-price")
Emitter<String> emitter;
```

Patterns

REST API → Message

```
@Inject @Stream("in")
Emitter<String> emitter;

emitter.send(message);
```

Message → Message

```
@Incoming("in")
@Outgoing("out")
public String process(String in) {
}
```

Message → SSE

```
@Inject @Stream("out")
Publisher<String> result;

@GET
@Produces(SERVER_SENT_EVENTS)
public Publisher<String> stream() {
    return result;
}
```

Message → Business Logic

```
@ApplicationScoped
public class ReceiverMessages {
    @Incoming("prices")
    public void print(String price) {
    }
}
```

Possible implementations are:

In-Memory

If the stream is not configured then it is assumed to be an in-memory stream, if not then stream type is defined by `connector` field.

Kafka

To integrate with Kafka you need to add next extensions:

```
mvn quarkus:add-extension
    -Dextensions="
    io.quarkus:quarkus-smallrye-reactive-messaging-kafka"
```

Then `@Outgoing`, `@Incoming` or `@Stream` can be used.

Kafka configuration schema: `mp.messaging.[outgoing|incoming].{stream-name}.<property>=<value>`.

The `connector` type is `smallrye-kafka`.

```
mp.messaging.outgoing.generated-price.connector=
    smallrye-kafka
mp.messaging.outgoing.generated-price.topic=
    prices
mp.messaging.outgoing.generated-price.bootstrap.servers=
    localhost:9092
mp.messaging.outgoing.generated-price.value.serializer=
    org.apache.kafka.common.serialization.IntegerSerializer

mp.messaging.incoming.prices.connector=
    smallrye-kafka
mp.messaging.incoming.prices.value.deserializer=
    org.apache.kafka.common.serialization.IntegerDeserializer
```

A complete list of supported properties are in Kafka site. For the producer and for consumer

JSON-B Serializer/Deserializer

You can use JSON-B to serialize/deserialize objects.

```
./mvnw quarkus:add-extension
    -Dextensions="quarkus-kafka-client"
```

To serialize you can use `io.quarkus.kafka.client.serialization.JsonbSerializer`.

To deserialize you need to extend `io.quarkus.kafka.client.serialization.JsonbDeserializer` and provide a type.

```
public class BeerDeserializer
    extends JsonbDeserializer<Beer> {

    public BeerDeserializer() {
        super(Beer.class);
    }

}
```

AMQP

To integrate with AMQP you need to add next extensions:

```
./mvnw quarkus:add-extension
    -Dextensions="reactive-messaging-amqp"
```

Then `@Outgoing`, `@Incoming` or `@Stream` can be used.

AMQP configuration schema: `mp.messaging.[outgoing|incoming].{stream-name}.<property>=<value>`. Special properties `amqp-username` and `amqp-password` are used to configure AMQP broker credentials.

The connector type is `smallrye-amqp`.

```
amqp-username=quarkus
amqp-password=quarkus
# write
mp.messaging.outgoing.generated-price.connector=
    smallrye-amqp
mp.messaging.outgoing.generated-price.address=
    prices
mp.messaging.outgoing.generated-price.durable=
    true
# read
mp.messaging.incoming.prices.connector=
    smallrye-amqp
mp.messaging.incoming.prices.durable=
    true
```

A complete list of supported properties for AMQP.

MQTT

To integrate with MQTT you need to add next extensions:

```
./mvnw quarkus:add-extension
    -Dextensions="vertx, smallrye-reactive-streams-operators,
        smallrye-reactive-messaging"
```

And add `io.smallrye.reactive:smallrye-reactive-messaging-mqtt-1.0:0.0.10` dependency in your build tool.

Then `@Outgoing`, `@Incoming` or `@Stream` can be used.

MQTT configuration schema: `mp.messaging.[outgoing|incoming].{stream-name}.<property>=<value>`.

The connector type is `smallrye-mqtt`.

```
mp.messaging.outgoing.topic-price.type=
    smallrye-mqtt
mp.messaging.outgoing.topic-price.topic=
    prices
mp.messaging.outgoing.topic-price.host=
    localhost
mp.messaging.outgoing.topic-price.port=
    1883
mp.messaging.outgoing.topic-price.auto-generated-client-id=
    true

mp.messaging.incoming.prices.type=
    smallrye-mqtt
mp.messaging.incoming.prices.topic=
    prices
mp.messaging.incoming.prices.host=
    localhost
mp.messaging.incoming.prices.port=
    1883
mp.messaging.incoming.prices.auto-generated-client-id=
    true
```

Kafka Streams

Create streaming queries with the Kafka Streams API.

```
./mvnw quarkus:add-extension
    -Dextensions="kafka-streams"
```

You need to initialize `org.apache.kafka.streams.KafkaStreams` on startup:

```
private KafkaStreams streams;

void onStart(@Observes StartupEvent ev) {
    Properties props = new Properties();
    // ...
    StreamsBuilder builder = new StreamsBuilder();
    // ...
    streams = new KafkaStreams(builder.build(), props);

    executor = Executors.newSingleThreadExecutor();
    executor.execute(() -> {
        waitForTopicsToBeCreated(bootstrapServers);
        streams.start();
    });
}

void onStop(@Observes ShutdownEvent ev) {
    streams.close();
    executor.shutdown();
}
```

And then you can use it.

```
public List getData() {
    return streams.allMetadataForStore("")
                .stream()
                .map()
                .collect()
}
```

Reactive PostgreSQL Client

You can use Reactive PostgreSQL to execute queries to PostreSQL database in a reactive way, instead of using JDBC way.

```
./mvnw quarkus:add-extension
   -Dextensions="io.quarkus:quarkus-reactive-pg-client"
```

Database configuration is the same as shown in **Persistence** section, but URL is different as it is not a *jdbc*.

```
quarkus.datasource.url=
    vertx-reactive:postgresql://host:5431/db
```

Then you can inject `io.reactiverse.axle.pgclient.PgPool` class.

```
@Inject
PgPool client;

CompletionStage<JsonArray> =
    client.query("SELECT * FROM table")
    .thenApply(pgRowSet -> {
        JsonArray jsonArray = new JsonArray();
        PgIterator iterator = pgRowSet.iterator();
        return jsonArray;
    })
```

JWT

Quarkus implements **MicroProfile JWT RBAC spec.**

```
mvn quarkus:add-extension
    -Dextensions="io.quarkus:quarkus-smallrye-jwt"
```

Minimum JWT required claims: `typ`, `alg`, `kid`, `iss`, `sub`, `exp`, `iat`, `jti`, `upn`, `groups`.

You can inject token by using `JsonWebToken` or a claim individually by using `@Claim`.

```
@Inject
JsonWebToken jwt;

@Inject
@Claim(standard = Claims.preferred_username)
String name;

@Inject
@Claim("groups")
Set<String> groups;
```

Set of supported types: `String`, `Set<String>`, `Long`, `Boolean`, `javax.json.JsonValue`, `Optional`, `org.eclipse.microprofile.jwt.ClaimValue`.

And configuration in `src/main/resources/application.properties`:

```
mp.jwt.verify.publickey.location=
    META-INF/resources/publicKey.pem
mp.jwt.verify.issuer=
    https://quarkus.io/using-jwt-rbac
```

Configuration options:

Parameter	Default	Description
`quarkus.smallrye-jwt.enabled`	true	Determine if the jwt extension is enabled.
`quarkus.smallrye-jwt.realm-name`	Quarkus-JWT	Name to use for security realm.
`quarkus.smallrye-jwt.auth-mechanism`	MP-JWT	Authentication mechanism.

Parameter	Default	Description
`mp.jwt.verify.publickey`	none	Public Key text itself to be supplied as a string.
`mp.jwt.verify.publickey.location`	none	Relative path or URL of a public key.
`mp.jwt.verify.issuer`	none	`iss` accepted as valid.

Supported public key formats:

- PKCS#8 PEM
- JWK
- JWKS
- JWK Base64 URL
- JWKS Base64 URL

To send a token to server-side you should use `Authorization` header: `curl -H "Authorization: Bearer eyJraWQiOi…"`.

To inject claim values, the bean must be `@RequestScoped` CDI scoped. If you need to inject claim values in scope with a lifetime greater than `@RequestScoped` then you need to use `javax.enterprise.inject.Instance` interface.

```
@Inject
@Claim(standard = Claims.iat)
private Instance<Long> providerIAT;
```

RBAC

JWT `groups` claim is directly mapped to roles to be used in security annotations.

```
@RolesAllowed("Subscriber")
```

Keycloak

Quarkus can use **Keycloak** to protect resources using bearer token issued by Keycloak server.

```
mvn quarkus:add-extension
    -Dextensions="io.quarkus:quarkus-keycloak"
```

You can get token information by injecting `KeycloakSecurityContext` object.

```
@Inject
KeycloakSecurityContext keycloakSecurityContext;
```

You can also protect resources with security annotations.

```
@GET
@RolesAllowed("admin")
```

Configure application to Keycloak service in `application.properties` file.

```
quarkus.keycloak.realm=quarkus
quarkus.keycloak.auth-server-url=http://localhost:8180/auth
quarkus.keycloak.resource=backend-service
quarkus.keycloak.bearer-only=true
quarkus.keycloak.credentials.secret=secret
quarkus.keycloak.policy-enforcer.enable=true
quarkus.keycloak.policy-enforcer.enforcement-mode=PERMISSIVE
quarkus.keycloak.enable-cors=true
quarkus.keycloak.cors-max-age=1000
quarkus.keycloak.cors-allowed-methods=POST, PUT, DELETE, GET
quarkus.keycloak.cors-exposed-headers=WWW-Authenticate
```

You can see all possible **Configuration parameters here**.

Tip you can also use `src/main/resources/keycloak.json` standard Keycloak configuration file.

OAuth2

Quarkus integrates with OAuth2 to be used in case of opaque tokens (none JWT) and its validation against an introspection endpoint.

```
mvn quarkus:add-extension
    -Dextensions="security-oauth2"
```

And configuration in `src/main/resources/application.properties`:

```
quarkus.oauth2.client-id=client_id
quarkus.oauth2.client-secret=secret
quarkus.oauth2.introspection-url=http://oauth-server/introspect
```

And you can map roles to be used in security annotations.

```
@RolesAllowed("Subscriber")
```

Configuration options:

Parameter	Default	Description
`quarkus.oauth2.enabled`	true	Determine if the OAuth2 extension is enabled.
`quarkus.oauth2.client-id`		The OAuth2 client id used to validate the token.
`quarkus.oauth2.client-secret`		The OAuth2 client secret used to validate the token.
`quarkus.oauth2.introspection-url`		URL used to validate the token and gather the authentication claims.
`quarkus.oauth2.role-claim`	scope	The claim that is used in the endpoint response to load the roles

JAX-RS

Quarkus uses JAX-RS to define REST-ful web APIs.

```
@Path("/book")
public class BookResource {

    @GET
    @Produces(MediaType.APPLICATION_JSON)
    public List<Book> getAllBooks() {}

    @POST
    @Produces(MediaType.APPLICATION_JSON)
    public Response createBook(Book book) {}

    @DELETE
    @Path("{isbn}")
    @Produces(MediaType.APPLICATION_JSON)
    public Response deleteBook(
        @PathParam("isbn") String isbn) {}


    @GET
    @Produces(MediaType.APPLICATION_JSON)
    @Path("search")
    public Response searchBook(
        @QueryParam("description") String description) {}
}
```

To get information from request:

Property	Description Example	
@PathParam	Gets content from request URI.	/book/{id} @PathParam("id")
@QueryParam	Gets query parameter.	/book?desc="" @QueryParam("desc")
@FormParam	Gets form parameter.	
@MatrixParam	Get URI matrix parameter.	/book;author=mkyong;country=malaysia
@CookieParam	Gets cookie param by name.	
@HeaderParam	Gets header parameter by name.	

Valid HTTP method annotations provided by the spec are: `@GET`, `@POST`, `@PUT`, `@DELETE`, `@PATCH`, `@HEAD` and `@OPTIONS`.

You can create new annotations that bind to HTTP methods not defined by the spec.

```
@Target({ElementType.METHOD})
@Retention(RetentionPolicy.RUNTIME)
@HttpMethod("LOCK")
public @interface LOCK {
}

@LOCK
public void lockIt() {}
}
```

Injecting

Using `@Context` annotation to inject JAX-RS and Servlet information.

```
@GET
public String getBase(@Context UriInfo uriInfo) {
    return uriInfo.getBaseUri();
}
```

Possible injectable objects: `SecurityContext`, `Request`, `Application`, `Configuration`, `Providers`, `ResourceContext`, `ServletConfig`, `ServletContext`, `HttpServletRequest`, `HttpServletResponse`, `HttpHeaders`, `Urinfo`, `SseEventSink` and `Sse`.

HTTP Filters

HTTP request and response can be intercepted to manipulate the metadata (ie headers, parameters, media type, …) or abort a request. You only need to implement the next `ContainerRequestFilter` and `ContainerResponseFilter` JAX-RS interfaces respectively.

```
@Provider
public class LoggingFilter
        implements ContainerRequestFilter {

    @Context
    UriInfo info;

    @Context
    HttpServletRequest request;

    @Override
    public void filter(ContainerRequestContext context) {
        final String method = context.getMethod();
        final String path = info.getPath();
        final String address = request.getRemoteAddr();
        System.out.println("Request %s %s from IP %s",
                method, path, address);
    }

}
```

Exception Mapper You can map exceptions to produce a custom output by implementing `ExceptionMapper` interface:

```
@Provider
public class ErrorMapper
    implements ExceptionMapper<Exception> {

    @Override
    public Response toResponse(Exception exception) {
        int code = 500;
        if (exception instanceof WebApplicationException) {
            code = ((WebApplicationException) exception)
                .getResponse().getStatus();
        }
        return Response.status(code)
            .entity(
                Json.createObjectBuilder()
                .add("error", exception.getMessage())
                .add("code", code)
                .build()
            )
            .build();
    }
}
```

CORS Filter

Quarkus comes with a CORS filter that can be enabled via configuration:

```
quarkus.http.cors=true
```

Prefix is `quarkus.http`.

Property	Default	Description

Property	Default	Description
cors	false	Enable CORS.
cors.origins	Any request valid.	CSV of origins allowed.
cors.methods	Any method valid.	CSV of methods valid.
cors.headers	Any requested header valid.	CSV of valid allowed headers.
cors.exposed-headers		CSV of valid exposed headers.

Fault Tolerance

Quarkus uses MicroProfile Fault Tolerance spec:

```
./mvnw quarkus:add-extension
  -Dextensions="io.quarkus:quarkus-smallrye-fault-tolerance"
```

MicroProfile Fault Tolerance spec uses CDI interceptor and it can be used in several elements such as CDI bean, JAX-RS resource or MicroProfile Rest Client.

To do automatic **retries** on a method:

```
@Path("/api")
@RegisterRestClient
public interface WorldClockService {
    @GET @Path("/json/cet/now")
    @Produces(MediaType.APPLICATION_JSON)
    @Retry(maxRetries = 2)
    WorldClock getNow();
}
```

You can set fallback code in case of an error by using `@Fallback` annotation:

```
@Retry(maxRetries = 1)
@Fallback(fallbackMethod = "fallbackMethod")
WorldClock getNow(){}

public String fallbackMethod() {
    return "It could beworse";
}
```

`fallbackMethod` must have the same parameters and return type as the annotated method.

You can also set logic into a class that implements `FallbackHandler` interface:

```
public class RecoverFallback
            implements FallbackHandler<String> {
    @Override
    public String handle(ExecutionContext context) {
    }
}
```

And set it in the annotation as value `@Fallback(RecoverFallback.class)`.

In case you want to use **circuit breaker** pattern:

```
@CircuitBreaker(requestVolumeThreshold = 4,
                failureRatio=0.75,
                delay = 1000)
WorldClock getNow(){}
```

If 3 (`4 x 0.75`) failures occur among the rolling window of 4 consecutive invocations then the circuit is opened for 1000 ms and then be back to half open. If the invocation succeeds then the circuit is back to closed again.

You can use **bulkhead** pattern to limit the number of concurrent access to the same resource. If the operation is synchronous it uses a semaphore approach, if it is asynchronous a thread-pool one. When a request cannot be processed `BulkheadException` is thrown. It can be used together with any other fault tolerance annotation.

```
@Bulkhead(5)
@Retry(maxRetries = 4,
       delay = 1000,
       retryOn = BulkheadException.class)
WorldClock getNow(){}
```

Fault tolerance annotations:

Annotation	Properties
@Timeout	unit
@Retry	maxRetries, delay, delayUnit, maxDuration, durationUnit, jitter, jitterDelayUnit, retryOn, abortOn
@Fallback	fallbackMethod
@Bulkhead	waitingTaskQueue (only valid in asynchronous)

Annotation	Properties
@CircuitBreaker	failOn, delay, delayUnit, requestVolumeThreshold, failureRatio, successThreshold
@Asynchronous	

You can override annotation parameters via configuration file using property `[classname/methodname/]annotation/parameter`:

```
org.acme.quickstart.WorldClock/getNow/Retry/maxDuration=30
# Class scope
org.acme.quickstart.WorldClock/Retry/maxDuration=3000
# Global
Retry/maxDuration=3000
```

You can also enable/disable policies using special parameter `enabled`.

```
org.acme.quickstart.WorldClock/getNow/Retry/enabled=false
# Disable everything except fallback
MP_Fault_Tolerance_NonFallback_Enabled=false
```

Tip MicroProfile Fault Tolerance integrats with MicroProfile Metrics spec. You can disable it by setting `MP_Fault_Tolerance_Metrics_Enabled` to false.

Observability

Health Checks

Quarkus relies on MicroProfile Health spec to provide health checks.

```
./mvnw quarkus:add-extension
  -Dextensions="io.quarkus:quarkus-smallrye-health"
```

By just adding this extension, an endpoint is registered to `/health` providing a default health check.

```
{
    "status": "UP",
    "checks": [
    ]
}
```

To create a custom health check you need to implement the `HealthCheck` interface and annotate either with `@Readiness` (ready to process requests) or `@Liveness` (is running)

```
@Readiness
public class DatabaseHealthCheck implements HealthCheck {
    @Override
    public HealthCheckResponse call() {
        HealthCheckResponseBuilder responseBuilder =
            HealthCheckResponse.named("Database conn");

        try {
            checkDatabaseConnection();
            responseBuilder.withData("connection", true);
            responseBuilder.up();
        } catch (IOException e) {
            // cannot access the database
            responseBuilder.down()
                .withData("error", e.getMessage());
        }
        return responseBuilder.build();
    }
}
```

Builds the next output:

```
{
    "status": "UP",
    "checks": [
        {
            "name": "Database conn",
            "status": "UP",
            "data": {
                "connection": true
            }
        }
    ]
}
```

Since health checks are CDI beans, you can do:

```
@ApplicationScoped
public class DatabaseHealthCheck {

    @Produces
    @ApplicationScoped
    @Liveness
    HealthCheck check1() {
      return io.smallrye.health.HealthStatus
                            .up("successful-live");
    }

    @Produces
    @ApplicationScoped
    @Readiness
    HealthCheck check2() {
      return HealthStatus
              .state("successful-read", this::isReady)
    }

    private boolean isReady() {}
}
```

You can ping liveness or readiness health checks individually by querying `/health/live` or `/health/ready`.

Metrics

Quarkus can utilize the MicroProfile Metrics spec to provide metrics support.

```
./mvnw quarkus:add-extension
    -Dextensions="io.quarkus:quarkus-smallrye-metrics"
```

The metrics can be read with JSON or the OpenMetrics format. An endpoint is registered automatically at `/metrics` providing default metrics.

MicroProfile Metrics annotations:

Annotation	Description
`@Timed`	Tracks the duration.
`@Metered`	Tracks the frequency of invocations.
`@Counted`	Counts number of invocations.
`@Gauge`	Samples the value of the annotated object.
`@ConcurrentGauge`	Gauge to count parallel invocations.
`@Metric`	Used to inject a metric. Valid types `Meter`, `Timer`, `Counter`, `Histogram`. `Gauge` only on producer methods/fields.

```
@GET
//...
@Timed(name = "checksTimer",
description = "A measure of how long it takes
                              to perform a hello.",
unit = MetricUnits.MILLISECONDS)
public String hello() {}

@Counted(name = "countWelcome",
description = "How many welcome have been performed.")
public String hello() {}
```

`@Gauge` annotation returning a measure as a gauge.

```
@Gauge(name = "hottestSauce", unit = MetricUnits.NONE,
description = "Hottest Sauce so far.")
```

Injecting a histogram using `@Metric`.

```
@Inject
@Metric(name = "histogram")
Histogram historgram;
```

Tracing

Quarkus can utilize the MicroProfile OpenTracing spec.

```
./mvnw quarkus:add-extension
    -Dextensions="io.quarkus:quarkus-smallrye-opentracing"
```

Requests sent to any endpoint are traced automatically.

This extension includes OpenTracing support and `Jaeger` tracer.

Jaeger tracer configuration:

```
quarkus.jaeger.service-name=myservice
quarkus.jaeger.sampler-type=const
quarkus.jaeger.sampler-param=1
quarkus.jaeger.endpoint=http://localhost:14268/api/traces
```

`@Traced` annotation can be set to disable tracing at class or method level.

`Tracer` class can be injected into the class.

```
@Inject
Tracer tracer;

tracer.activeSpan().setBaggageItem("key", "value");
```

Cloud

Native

You can build a native image by using GraalVM. The common use case is creating a Docker image so you can execute the next commands:

```
./mvnw package -Pnative -Dnative-image.docker-build=true

docker build -f src/main/docker/Dockerfile.native
                -t quarkus/getting-started .
docker run -i --rm -p 8080:8080 quarkus/getting-started
```

To configure native application, you can create a `config` directory at the same place as the native file and place an `application.properties` file inside. `config/application.properties`.

Kubernetes

Quarks can use Dekorate to generate Kubernetes resources.

```
./mvnw quarkus:add-extensions
    -Dextensions="io.quarkus:quarkus-kubernetes"
```

Running `./mvnw package` the Kubernetes resources are created at `target/wiring-classes/META-INF/kubernetes/` directory.

Property	Default	Description
quarkus.kubernetes.group	Current username	Set Docker Username.
quarkus.application.name	Current project name	Project name

Generated resource is integrated with MicroProfile Health annotations.

Kubernetes Client

Quarkus integrates with Fabric8 Kubernetes Client.

```
./mvnw quarkus:add-extension
    -Dextensions="quarkus-kubernetes-client"
```

List of Kubernetes client parameters.

`quarkus.kubernetes-client` as prefix is skipped in the next table.

Property	Default	Description

Property	Default	Description
trust-certs	false	Trust self-signed certificates.
master-url		URL of Kubernetes API server.
namesapce		Default namespace.
ca-cert-file		CA certificate data.
client-cert-file		Client certificate file.
client-cert-data		Client certificate data.
client-key-data		Client key data.
client-key-algorithm		Client key algorithm.
client-key-passphrase		Client key passphrase.
username		Username.
password		Password.
watch-reconnect-interval	PT1S	Watch reconnect interval.
watch-reconnect-limit	-1	Maximum reconnect attempts.
connection-timeout	PT10S	Maximum amount of time to wait for a connection.
request-timeout	PT10S	Maximum amount of time to wait for a request.
rolling-timeout	PT15M	Maximum amount of time to wait for a rollout.

Property	Default	Description
http-proxy		HTTP proxy used to access the Kubernetes.
https-proxy	``	HTTPS proxy used to access the Kubernetes.
proxy-username		Proxy username.
proxy-password		Proxy password.
no-proxy		IP addresses or hosts to exclude from proxying

Or programmatically:

```
@Dependent
public class KubernetesClientProducer {

    @Produces
    public KubernetesClient kubernetesClient() {
        Config config = new ConfigBuilder()
            .withMasterUrl("https://mymaster.com")
            .build();
        return new DefaultKubernetesClient(config);
    }
}
```

And inject it on code:

```
@Inject
KubernetesClient client;

ServiceList myServices = client.services().list();

Service myservice = client.services()
        .inNamespace("default")
        .withName("myservice")
        .get();

CustomResourceDefinitionList crds = client
        .customResourceDefinitions()
        .list();

dummyCRD = new CustomResourceDefinitionBuilder()
        ...
        .build()
client.customResourceDefinitions()
        .create(dummyCRD);
```

Testing Quarkus provides a Kubernetes Mock test resource that starts a mock of Kubernetes API server and sets the proper environment variables needed by Kubernetes Client.

Register next dependency: `io.quarkus:quarkus-test-kubernetes-client:test`.

```java
@QuarkusTestResource(KubernetesMockServerTestResource.class)
@QuarkusTest
public class KubernetesClientTest {

    @MockServer
    private KubernetesMockServer mockServer;

    @Test
    public void test() {
        final Pod pod1 = ...
        mockServer
            .expect()
            .get()
            .withPath("/api/v1/namespaces/test/pods")
            .andReturn(200,
                new PodListBuilder()
                    .withNewMetadata()
                    .withResourceVersion("1")
                    .endMetadata()
                    .withItems(pod1, pod2)
                    .build())
            .always();
    }
}
```

Amazon Lambda

Quarkus integrates with Amazon Lambda.

```
./mvnw quarkus:add-extension
  -Dextensions="io.quarkus:quarkus-amazon-lambda"
```

And then implement `com.amazonaws.services.lambda.runtime.RequestHandler` interface.

```java
public class TestLambda
        implements RequestHandler<MyInput, MyOutput> {
    @Override
    public MyInput handleRequest(MyOutput input,
                                 Context context) {

    }
}
```

Test

You can write tests for Amazon lambdas:

```xml
<dependency>
  <groupId>io.quarkus</groupId>
  <artifactId>quarkus-test-amazon-lambda</artifactId>
  <scope>test</scope>
</dependency>
```

```java
@Test
public void testLambda() {
    MyInput in = new MyInput();
    in.setGreeting("Hello");
    in.setName("Stu");
    MyOutput out = LambdaClient.invoke(MyOutput.class, in);
}
```

Apache Camel

Quarkus integrates wih **Apache Camel**.

```
./mvnw quarkus:add-extension
  -Dextensions="io.quarkus:quarkus-camel-core"
```

You can inject `CamelRuntime` and `CamelContext` instances:

```java
@Inject
CamelRuntime runtime;

@Inject
CamelContext context;
```

And also observing Camel events `InitializingEvent`, `InitializedEvent`, `StartingEvent`, `StartedEvent`. `StoppingEvent`, `StoppedEvent`.

```java
public void init(@Observes InitializingEvent event) {
    runtime.getContext()
            .setUuidGenerator(new DefaultUuidGenerator());
}
```

You can create `RouteBuilder` for routing rules:

```java
public class CamelRoute extends RouteBuilder {

    @Override
    public void configure() {
        from("timer:keep-alive?" +
                "period={{camel.timer-route.period}}")
        .id("timer-route")
        .setBody(constant(""))
        .to("log:keep-alive");
    }
}
```

Custom properties are set in `application.properties`.

```
camel.timer-route.period=1s
```

List of Camel parameters.

`quarkus.` as a prefix is skipped in the next table.

Parameter	Default	Description
camel.routesUris		Uri to an XML with camel routes.

Parameter	Default	Description
`camel.deferInitPhase`	false	Defer context initialization phase until runtime.
`camel.disableJaxb`	false	Disable jaxb support at runtime. Useful when routes loaded at build time.
`camel.disableXml`	false	Disable XML support in various parts of Camel.
`camel.dumpRoutes`	false	Dump loaded routes .

Quarkus also comes with support for Camel in form of next extensions: `camel-aws-s3`, `camel-aws-sqs`, `camel-infinispan`, `camel-netty4-http`, `camel-salesforce`.

WebSockets

Quarkus can be used to handling web sockets.

```
./mvnw quarkus:add-extension
  -Dextensions="io.quarkus:quarkus-undertow-websockets"
```

And web sockets classes can be used:

```
@ServerEndpoint("/chat/{username}")
@ApplicationScoped
public class ChatSocket {

    @OnOpen
    public void onOpen(Session session,
               @PathParam("username") String username) {}

    @OnClose
    public void onClose(..) {}

    @OnError
    public void onError(..., Throwable throwable) {}

    @OnMessage
    public void onMessage(...) {}

}
```

OpenAPI

Quarkus can expose its API description as OpenAPI spec and test it using Swagger UI.

```
./mvnw quarkus:add-extension
  -Dextensions="io.quarkus:quarkus-smallrye-openapi"
```

Then you only need to access to `/openapi` to get OpenAPI v3 spec of services.

You can update the OpenApi path by setting `quarkus.smallrye-openapi.path` property.

Also, in case of starting Quarkus application in `dev` or `test` mode, Swagger UI is accessible at `/swagger-ui`. If you want to use it in production mode you need to set `quarkus.swagger-ui.always-include` property to `true`.

You can update the Swagger UI path by setting `quarkus.swagger-ui.path` property.

```
quarkus.swagger-ui.path=/my-custom-path
```

You can customize the output by using Open API v3 annotations.

```
@Schema(name="Developers", description="POJO that represents a developer.")
public class Developer {
    @Schema(required = true, example = "Alex")
    private String name;
}

@POST
@Path("/developer")
@Operation(summary = "Create deeloper",
          description = "Only be done by admin.")
public Response createDeveloper(
        @RequestBody(description = "Developer object",
            required = true,
            content = @Content(schema =
            @Schema(implementation = Developer.class)))
                Developer developer)
```

All possible annotations can be seen at org.eclipse.microprofile.openapi.annotations package.

Mail Sender

You can send emails by using Quarkus Mailer extension:

```
./mvnw quarkus:add-extension
  -Dextensions="io.quarkus:quarkus-mailer"
```

You can inject two possible classes `io.quarkus.mailer.Mailer` for synchronous API or `io.quarkus.mailer.ReactiveMailer` for asynchronous API.

```
@Inject
Mailer mailer;

@Inject
ReactiveMailer reactiveMailer;
```

And then you can use them to send an email:

```
mailer.send(
    Mail.withText("to@acme.org", "Subject", "Body")
);

CompletionStage<Void> stage =
    reactiveMailer.send(
        Mail.withText("to@acme.org", "Subject", "Body")
    );
```

`Mail` class contains methods to add `cc`, `bcc`, `headers`, `bounce address`, `reply to`, `attachments`, `inline attachments` and `html body`.

```
mailer.send(Mail.withHtml("to@acme.org", "Subject", body)
        .addInlineAttachment("quarkus.png",
            new File("quarkus.png"),
            "image/png", "<my-image@quarkus.io>"));
```

Tip If you need deep control you can inject Vert.x mail client `@Inject MailClient client;`

You need to configure SMTP properties to be able to send an email:

```
quarkus.mailer.from=test@quarkus.io
quarkus.mailer.host=smtp.sendgrid.net
quarkus.mailer.port=465
quarkus.mailer.ssl=true
quarkus.mailer.username=....
quarkus.mailer.password=....
```

List of Mailer parameters. `quarkus.` as a prefix is skipped in the next table.

Parameter	Default	Description
`mailer.from`		Default address.

Parameter	Default	Description
`mailer.mock`	false in `prod`, true in `dev` and `test`.	Emails not sent, just printed and stored in a `MockMailbox` bean.
`mailer.bounce-address`		Default address.
`mailer.host`	*mandatory*	SMTP host.
`mailer.port`	25	SMTP port.
`mailer.username`		The username.
`mailer.password`		The password.
`mailer.ssl`	`false`	Enables SSL.
`mailer.trust-all`	false	Trust all certificates.
`mailer.max-pool-size`	10	Maximum number of open connections .
`mailer.own-host-name`		Hostname for `HELO/EHLO` and `Message-ID`
`mailer.keep-alive`	`true`	Connection pool enabled.
`mailer.disable-esmtp`	`false`	Disable ESMTP.
`mailer.start-tls`	`OPTIONAL`	TLS security mode. `NONE`, `OPTIONAL`, `REQUIRED`.
`mailer.login`	`NONE`	Login mode. `NONE`, `OPTIONAL`, `REQUIRED`.

Parameter	Default	Description
`mailer.auth-methods`	All methods.	Space-separated list.
`mailer.key-store`		Path of the key store.
`mailer.key-store-password`		Key store password.

IMPORTANT: if you enable SSL for the mailer and you want to build a native executable, you will need to enable the SSL support `quarkus.ssl.native=true`.

Testing

If `quarkus.mailer.mock` is set to `true`, which is the default value in `dev` and `test` mode, you can inject `MockMailbox` to get the sent messages.

```
@Inject
MockMailbox mailbox;

@BeforeEach
void init() {
    mailbox.clear();
}

List<Mail> sent = mailbox
                    .getMessagesSentTo("to@acme.org");
```

Scheduled Tasks

You can schedule periodic tasks with Quarkus.

```
@ApplicationScoped
public class CounterBean {

    @Scheduled(every="10s")
    void increment() {}

    @Scheduled(cron="0 15 10 * * ?")
    void morningTask() {}
}
```

`every` and `cron` parameters can be surrounded with `{}` and the value is used as config property to get the value.

```
@Scheduled(cron = "{morning.check.cron.expr}")
void morningTask() {}
```

And configure the property into `application.properties`:

```
morning.check.cron.expr=0 15 10 * * ?
```

Kogito

Quarkus integrates with **Kogito**, a next-generation business automation toolkit from Drools and jBPM projects for adding business automation capabilities.

To start using it you only need to add the next extension:

```
./mvnw quarkus:add-extension
  -Dextensions="kogito"
```

Apache Tika

Quarkus integrets with **Apache Tika** to detect and extract metadata/text from different file types:

```
./mvnw quarkus:add-extension
  -Dextensions="quarkus-tika"
```

```
@Inject
io.quarkus.tika.TikaParser parser;

@POST
@Path("/text")
@Consumes({ "text/plain", "application/pdf",
            "application/vnd.oasis.opendocument.text" })
@Produces(MediaType.TEXT_PLAIN)
public String extractText(InputStream stream) {
    return parser.parse(stream).getText();
}
```

Resources

- https://quarkus.io/guides/
- https://www.youtube.com/user/lordofthejars

Authors :

@alexsotob
Java Champion and SW Engineer at Red Hat